JULIAN BHARADWAJA

Tustin, CA | bharadwajajulian@gmail.com
linkedin.com/in/julian-bharadwaja | github.com/jbharadwaja | jbharadwaja.github.io/my-portfolio

SUMMARY

Software engineer retraining in machine learning. Did DevOps at Fujitsu in Tokyo, then full-stack development for a South Korean company. Builds ML models in Python and reports results as they come out, including a forecast where plain linear regression beat gradient boosting and a model that barely beat its median baseline. Wants to work on empirical AI safety research.

EDUCATION

Asian Institute of Technology, B.E. in Mechatronics Engineering, Thailand 2012-2017

Orange Coast College, A.S. in Computer Science, GPA 4.0 Expected Dec 2026

TripleTen, Data Science program Expected Dec 2026

- Coursework includes matrix algebra, gradient descent, neural networks, NLP and time series forecasting.

SKILLS

Python: pandas, NumPy, scikit-learn, LightGBM, CatBoost, spaCy, SciPy, Prophet, Matplotlib

Methods: classification and regression, gradient boosting, time series forecasting, text classification with TF-IDF, class imbalance, cross-validation, bootstrapping, hypothesis testing

Engineering: SQL, Git and GitHub, shell scripting, CI/CD pipelines, REST APIs, TypeScript, React, Vue.js, Laravel

MACHINE LEARNING PROJECTS

Notebooks and write-ups: github.com/jbharadwaja/Data_projects_TripleTen

Movie review sentiment classifier | Python, spaCy, scikit-learn

- Lemmatized IMDB reviews with spaCy, built TF-IDF features and trained logistic regression. Test F1 0.90 (target 0.85).

Used car price prediction | scikit-learn, LightGBM

- A code reviewer found data leakage. Refit imputation and encoders on training data only and added a held-out test set.
- Compared 6 model types on accuracy, training time and prediction speed. LightGBM was best, with test RMSE of 1,583 euros.

Gold recovery prediction | scikit-learn, cross-validation

- Predicted gold recovery at two purification stages, scored with a custom sMAPE metric and cross-validation.
- Test sMAPE 9.41, only slightly better than a median baseline (9.44). The README reports it as a near-tie.

Taxi demand forecasting | time series, scikit-learn, LightGBM, CatBoost

- Forecast next-hour airport taxi orders from lag and rolling-mean features. Plain linear regression beat LightGBM and CatBoost, with test RMSE 34.5 (target 48 or lower).

Oil well location under risk | scikit-learn, bootstrapping

- Picked the best of 3 regions for 200 new wells, using linear regression for reserves and 1,000-sample bootstrapping for profit and risk of loss. Chosen region: $6.64M average profit, 0.1% risk of loss.

SHINYOROI, personal project | Python, PyTrends, Prophet

- Forecasts product interest from Google Trends data, with a backend and a mobile app.

ENGINEERING EXPERIENCE

Full-Stack Developer, Digital Strategic Solution, South Korea (hybrid) Aug 2021-Dec 2023

- Built and shipped an eCommerce platform in Vue.js and Laravel, including checkout, payment gateway integration and product search.
- Built an HR SaaS web app in TypeScript and React for employee data management and reporting.

Junior DevOps Engineer, Fujitsu, Tokyo, Japan Jan 2019-Sep 2020

- Automated build, test and deployment pipelines with shell scripts and CI/CD tooling.
- Monitored infrastructure through system logs and alerts, and helped move legacy components to a microservices architecture.

OTHER EXPERIENCE

Store Manager, The UPS Store, Tustin, CA Jan 2023-Present

- Own the store's revenue and cost reporting. Analyzed transaction and customer-mix data to find the highest-margin services, which contributed to a 25% increase in monthly revenue.
- Replaced manual reporting with weekly tracking of volume, revenue per transaction and service mix.

Student Experience & Program Insights Coordinator, Orange Coast College Jul 2023-Mar 2025

- Analyzed participation data across 20+ programs. The findings shaped program changes that led to a 30% rise in engagement.